

Step 2 — what the work is, and what to automate

Derived from `out/segments.jsonl` by `src/analyze.py` and `src/rank.py`. All figures are **relative**. The brief states the recordings were made in a test environment with compressed waiting times, so absolute durations are meaningless and only comparisons between processes are valid.

Scale

668 executions, 175 minutes of observed work, 15 sessions, 4 operators, all on 2026-07-01 between 16:44 and 19:36 UTC.

How many people, resolved

Three sources disagreed and had to be reconciled:

source	count
<code>username_hash</code>	4
<code>machine_id</code>	4 (strictly 1:1 with the hashes)
operator names on the portal dashboard	3

The three names are **not** operators — all three appear in every session. Each maps to one portal system:

name	system	consistency
田中 麻里・人事マネージャー	HR人事給与システム	98%
佐藤 直樹・経理マネージャー	財務会計システム	100%
高橋 健司・物流マネージャー	受発注在庫管理システム	97%

So: **4 people, working across all three systems under shared per-system logins**. That last point is a governance finding, not a trivia one — automation would run under a shared account, leaving no per-user audit trail.

The route names lie

The portal is one SPA deployed three times, so its route names repeat across all three systems and describe nothing. The regulation document open during the work says what the work actually is:

segment label	regulation document consulted	what it really is
ops__leave-applications	keiyaku_kaijo_tetsuzuki	contract termination
fin__leave-applications	settai_keihi_kitei	entertainment expense approval
fin__payroll-items	getsujitsu_teigaku_torihikisaki_ichiran	recurring supplier payment
fin__resident-tax	shinkuitorihikisaki_touroku_tetsuzuki	new supplier registration
hr__onboarding	nyusha_checklist_shinsotsu_batch	new-graduate onboarding
ops__social-insurance	kanrisha_kengen_shinsei_tetsuzuki	admin privilege request

This is also **independent validation of Step 1's labels**. The pipeline never reads document names, yet segments sharing a label consult the same document at **79% weighted purity**. That agreement is evidence the groupings track real business processes rather than portal geography.

Caveat: the English readings are mine and unverified — see [docs/glossary.md](#) .

Ranking

Two measurable quantities in tension:

- **Mechanical load** — clipboard events and application switches per run. The part automation removes.
- **Judgment load** — share of runs where a regulation document is open. The part automation does not remove, and what decides whether a deterministic script is even possible.

High volume with high judgment is a poor first target however expensive it looks, because the residue is what costs the time.

process	n	min	share	median s	mech	judgment	people
payroll_item_maintenance	115	24.5	14.0%	11	2.3	11%	4
recurring_supplier_payment	74	22.8	13.1%	15	3.5	55%	4
contract_termination	65	22.0	12.6%	18	6.3	79%	4
new_grad_onboarding	59	20.2	11.5%	19	6.3	83%	4
inventory_payroll_items	65	15.1	8.6%	11	3.7	6%	4
new_supplier_registration	67	13.8	7.9%	11	1.9	22%	4
leave_application_review	58	12.4	7.1%	10	1.9	0%	3
contractor_payment_setup	40	12.0	6.9%	20	5.5	60%	3
admin_privilege_request	45	10.3	5.9%	11	3.1	18%	4
entertainment_expense_approval	36	10.2	5.8%	16	5.1	67%	2
childcare_leave_handling	33	7.4	4.2%	13	3.3	52%	3

Robustness

A priority formula is easy to make say what you want, so the ranking was scored under six different weightings (time alone, count alone, time x automatability, mechanical x count, and two combined forms). Processes with $n < 5$ were excluded, since per-run rates on one or two observations are noise.

process	top-3 appearances
payroll_item_maintenance	6 / 6
recurring_supplier_payment	4 / 6
inventory_payroll_items	3 / 6
contract_termination	2 / 6
four others	1 / 6 each

Only one process survives every weighting. Anything appearing once is an artefact of a particular formula, not a finding.

The result that changes the recommendation

Aggregating by portal **screen** rather than by system:

screen	executions	minutes	share of all work	systems	judgment
payroll-items	254	62.4	35.7%	3	24%
leave-applications	159	44.6	25.5%	3	48%
onboarding	99	32.2	18.4%	2	72%
social-insurance	86	20.4	11.7%	3	23%
resident-tax	67	13.8	7.9%	1	22%

One screen pattern accounts for 36% of all observed work, occurs in all three systems, and carries the second-lowest judgment load. The three top-ranked processes are the same screen in three different deployments.

That reframes Step 3. The choice is not "which process do I automate" but "one bespoke process, or the shared pattern behind three of them" — which is exactly the scope question the brief says is itself part of the ROI decision.